Automated analysis and planning of online marketing campaigns

Erik Brännström
Chalmers University of Technology

1   Introduction

Duego was founded in 2010 with the idea of creating a social networking site targeting those who wish to meet new people. This stands in contrast to the majority of the existing networks that either manage already existing connections (e.g. Facebook, LinkedIn), or where the goal is the help the user to find a romantic partner (i.e. dating sites). Online advertisement is leveraged as a way to promote the site to new users. This is done both using targeted ads that are adapted based on the target demographics on social networks, as well as more conventional ads on websites. Today, managing these campaigns is a manual process that is outsourced to another company.

The goal of Duego is to automate the process of adapting ads based on target requirements using software. Such a system would be required to analyze existing campaigns with regard to ads and target groups along with campaign metrics, for example how many people clicked on an ad and how many times an ad has been shown. This data should then be used as a basis for suggesting new ads. The data set that will be used in production consists of hundreds of different campaigns, each with a large number of ads. Along with the attributes of both campaigns and ads, this means that the complete advertisement data set consists of hundreds of thousands of possible attribute combinations that need to be analyzed.

2   Marketing terminology

A number of marketing terms will be used throughout this paper. A *marketing campaign*, or simply a *campaign*, is comprised of one or more advertising messages, *ads*, that are directed to one defined audience, the *target* or *target group*. Any ad, and by extension campaign, have numeric measures of success called *metrics*. The most commonly used are *impressions*, which is the number of times an ad has been shown, and *clicks*. Ads are limited to include a title, a body text and/or an image, as well as the mandatory URL that a user who clicks the ad will be redirected to. A target is defined by demographic data such as age, gender, location or, in the case of for example search engine advertisement, search terms.

As the word user is typically used to refer to a person interacting with the web site that is being marketed, we will use the word *operator* when discussing
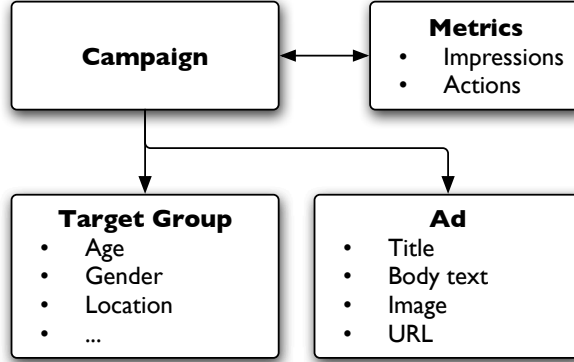
Figure 1: Marketing terminology

a person interacting with the system that is described in this paper to avoid confusion.

3 Problem

This problem statement is described from the view of Duego, however it can be generalized into a broader problem in the sphere of software engineering, namely knowledge discovery in databases (KDD) or more specifically data mining. Data mining is the process of delegating and automating the task of identifying useful knowledge in a large set of data to a computer, either fully or partially. The somewhat philosophical discussion of what knowledge really is and how to define what is useful knowledge will not however, be covered in this paper.


Figure 2: General MAPE control loop

Useful terminology is defined by IBM [2006] in the field of autonomic computing. Figure 2 shows an autonomic manager, which is a component that collects data from a system and, based on this data, performs actions with the purpose of improving the system. This control loop is divided into four subtasks called monitor (collect system information), analyze (correlate and model data),

plan (design behaviour required to reach goal) and execute (run the planned actions), sometimes refered to as MAPE. Each subtask can optionally interact with a knowledge base for storing and retrieving data.

In online marketing, monitoring is the collection of metrics for online advertisement. To optimize these metrics, the gathered information is analyzed and this analysis forms the basis for the planning of future campaigns. Once the plans are completed, the new campaign can be launched, with new metrics being gathered and so on.

To formalize the marketing terminology described previously, a target group is an m-tuple $T = (t_1, t_2, ..., t_m)$ and an ad is an n-tuple $A = (a_1, a_2, ..., a_n)$. A campaign can thus be described as the pair $C = (T, A)$ or, by flattening, $C' = (t_1, t_2, ..., t_m, a_1, a_2, ..., a_n)$. C' will typically be the format that the system expects input.
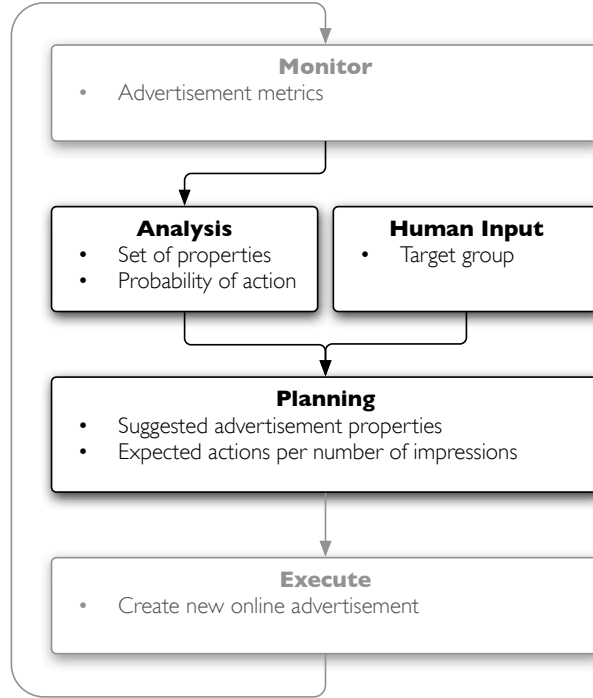

Figure 3: Online marketing automation system in control loop

Today only the monitoring is automated, where as the other parts of the process are performed manually. It is infeasible to fully automate the whole process, due to for example the creative side of advertisement and the complex factors that decide which groups an upcoming campaign should target. There are however certain areas that can be automated and this paper will focus on automated analysis and planning of online marketing campaigns, shown in Figure 3.

A system to automate this process would require monitored data as input, which in this context equals historical data of campaigns and their metrics as mentioned previously. This data would then be mined to identify subsets of

campaign properties that are associated with a probability that an impression of that an ad with these properties leads to an action being taken by the user. Human input is required to specify which target group the next campaign will be aimed at, and based on this the system will generate suggested ad properties that optimize the number of expected actions taken per impression.

4 Scope

Using the MAPE framework in Figure 2, only the analysis and planning tasks are considered part of this paper, where as monitoring and execution are out of scope. The latter two are however relevant in the verification step, but will not be covered as a research topic. In this context, this means for example that the feature of integrating this system with marketing services to automatically add new ads and campaigns will not be a part of the final system.

Furthermore, the data set includes attributes whose values are free text and images. Text mining and image recognition are beyond the scope of this project. Instead these attributes will be manually categorized, so that the value space is discrete and finite.

5 Foundations

A number of high-level descriptions of frameworks for knowledge discovery in databases exist [Fayyad et al., 1996, Frawley et al., 1992] and they exhibit a number of commonalities. These include the importance of having a knowledgeble human operator guiding the process in terms of supplying domain knowledge to the system formulating the goal of the knowledge discovery; feeding discovered knowledge back into the system; and the identification and application of a discovery method, or more specifically the data mining algorithm(s).

In data mining, the input to a system can be described using the terms *concepts*, *instances* and *attributes*, where concept is the actual result of the mining, i.e. what we want to be learned; an instance is one single example of data to be mined and can be compared to a row in a database; and attribute is a property of an instance, which in the database analogy would be a column [Witten et al., 2011].

The Weka Project has defined an input format to be used for their open source data mining software called ARFF (Attribute-Relation File Format) [Garner, 1995, Witten et al., 2011]. This format is used by the Weka software package, but is well-defined and can be used as the input format for custom systems as well. Witten et al. [2011] also describe how to use implementations of Weka in custom software projects as well as how to extend the system with new functionality.

In a highly influential paper, Quinlan [1986] decribe how decision trees can be created from a training set and how well it handles the problem of unknown attributes values and noisy data. A related paper, Quinlan [1987a], deal with how generated decision trees can be simplified in order to more easily be applied. Four different methods are evaluated, one of which is the reformulation of the tree as a set of production rules. This specific topic is further analyzed in Quinlan [1987b] where such production rules are shown to be more compact and

also in many cases improve the classification of unseen data. An added positive effect is that production rules from separate classifications can be merged more efficiently than their original decision trees.

Standard decision trees give a best-effort boolean classification of the input data, however sometimes it might be more appropriate to give the probability that the input belongs to each of the available classes. This can be described using probability estimation trees (PETs). Provost and Domingos [2003] discuss the problems of estimating these probabilities from ordinary decision trees and goes on to show how to increase the accuracy of the estimates by performing tree pruning more conservatively and by applying the Laplace correction. They also show that probability-bagging, meaning the combination of results from multiple classifiers instead of just the one, greatly improves the estimates. The suggested algorithm, called C4.4, is part of a comparative study of PETs by Chu et al. [2011] and is shown to have an impressive accuracy, though other algorithms may still be more appropriate. A likely candidate algorithm for this thesis is the Naive Bayes tree, which is a standard decision tree with the difference that the leaves are Naive Bayes classifiers, providing probabilities that an instance belong to each of the available classes.

A thorough literature study on domain-specific languages is covered by van Deursen et al. [2000] and covers high-level topics, such as potential benefits and risks of using DSLs and exemplifying with a number of languages from different areas. The article also mentions the phases of development in the creation of a new DSL, a topic which is described in more detail by Mernik et al. [2005]. The latter article also provides insight into the relevant choices that need be made in the creation process, and both articles help define a useful terminology for discussing the topic of domain-specific languages.

6 Related works

Chen et al. [1996] give an overview of the field of data mining where they mention the aspect of multi-level data mining, which states that correlations may not commonly exist on the lowest level of granularity, but instead by forming groups of related items. An example given would be that a specific brand of milk does not necessarily imply the purchase of a specific brand of bread, however purchasing milk of any kind may still be correlated to the purchase of bread irrespective of brand.

A related area in data mining is association rule learning. Agrawal et al. [1993] describe how, given a large set of transactions containing any number of different items, rules can be identified between these items. In the case of commerce, such rules would describe how the purchase of one product would, with a probability above a certain threshold, also infer the purchase of another product. This method could be applicable to the problem in this paper, even though it is more of a classification problem, and some of the necessary extensions of the method are mentioned below.

Srikant et al. [1997] expands the above research by applying contraints on the items in the resulting association rules. These constraints can either be expressed using specific items or with taxonomy rules, e.g. if an item is a descendant or ancestor of another item. For this problem, the relation between target groups and ads can, based on existing metrics, be restructured into a

transactional table where each transaction is an impression (in the case of the impression per click metrics) so that each entry would have an attribute stating whether or not the impression yielded a click. The item constraint would then be used to only formulate association rules that lead to clicks.

Considering that the original market basket problem assumes binary attributes and that relational data can contain both quantative and categorical data, [Srikant and Agrawal, 1996] gives a useful description on how such values can be mapped onto the binary (or boolean) problem space.

Hahsler et al. [2007] presents a package for the R software environment used in statistical computing, which implements a base for transaction databases as well as integration with two of the most common mining algorithms, Apriori and Eclat. One of the useful implementation choices is the implementation of the sparse matrix data structure, which greatly reduces memory load.

Richardson et al. [2007] describe a model for predicting the click-through rates of ads given a set of properties in the context of search engine advertisement. Though it is a relevant subject, the major difference between such advertisement and that which occurs on social networking sites is targeting. Search engines uses keyword matching to select possible ads, i.e. what the user is interested in at this moment, where as social networks can uitilize knowledge about who the user is.

7 Solution

The solution has been divided into one section each for the two subtasks identified in the problem description. Analysis covers the problem of creating an appropriate model of the available data that can be used in the next stage. Planning, in turn, is the utilization of that model with the purpose of outputting useful information to the operator.

7.1 Analysis

The natural distribution of the input data is very uneven because the number of impressions required to receive a click is high. Consider the scenario of a click-through rate (CTR) of 1%, which means that for every 100 impressions there would on average be one click. The simplest possible classification rule would be a rule that for every instance gives the classification of no click, and this would yield an accuracy of 99%. Even though this is indeed an impressive number, it is of no use to us since it says nothing about the instances that actually would yield clicks. This is refered to as the problem of imbalanced data sets as described by Chawla et al. [2004].

One way to deal with this problem is to increase the number of instances that we are interested in (positive instances) or decreasing the number of negative instances in so that the total distribution is even. This is called over-sampling and under- sampling respectively, and will remove the possibility for a classifier to simply disregard the positive instances in the classification [Chawla et al., 2004, Japkowicz and Stephen, 2002].

A better alternative is to use the concept of cost modification, i.e. instruct the algorithm that it is more costly to misclassify positive instances as negative instances than vice versa, which has been shown to be as effective as

over-sampling but with increased efficiency [Chawla et al., 2004, Japkowicz and Stephen, 2002] due to the fact that it does not increase the number of instances in the data set.

Once the monitored data has been passed to the system it will be added to the knowledge base. The data will also be used to generate a new analysis model of the full data set which also will be stored in the knowledge base, making it available for the planning subsystem.

As part of the generation of an analysis model, stored estimations may be compared to the monitored data and based on their differences the system can update a number of paramaters for the generation algorithm so that future estimates are improved.

7.1.1 Model generation

The model can be built by any data mining classifier algorithm that adheres to the basic requirements of handling both numerical and categorical values as well as generating either a decision tree or a set of decision rules, for example the PART algorithm [Frank and Witten, 1998].

To increase accuracy, a stratified tenfold cross-validation approach is recommended, which can optionally be repeated ten times to further improve the result with an additional cost in computing time. The validation step should calculate the success rate of the classification for each rule or leaf on each pass and average the values on completion. Confidence intervals can be calculated using tables for the Bernoulli distribution [Witten et al., 2011, chap. 5] based on the observed success rate S and the number of instances that matched the rule N. The probability of success is $\rho \pm z$ with confidence c.

7.1.2 TBR - Click probability

Once the appropriate cost matrix has been applied, the PART algorithm [Frank and Witten, 1998] is used to mine the data for rules. For each rule in the result, the number of instances in the data set that are covered by the rule (I) as well as how many of those instances that are misclassified (M) is calculated. The weight added to the positive instances is refered to as W. Using these known values, an expected probability of the rule producing a click, ρ, can be calculated.

$$\rho = \frac{(I-M)\frac{1}{W}}{(I-M)\frac{1}{W} + M} = \frac{I-M}{I-M+MW} = \frac{I-M}{I+M(W-1)}$$

This value is simply the number of correctly classified instances (i.e. actual clicks for impressions that match the rule) divided by the total number of matching instances (i.e. total number of impressions). W is used to convert the numbers back to their original ratio, since it is the factor used to increase the number of clicks as far as the algorithm is concerned.

7.2 Planning

The goal of the planning step is to identify relevant rules for new campaigns. The definition of relevant is created by the operator, who specifies rules for the target group and inputs them to the system. Based on these conditions of

relevancy together with the model created in the analysis step, relevant rules can be presented to the user.

Depending on the available training data, there might be targets that are not covered by the generated rule sets. By using rules that are similar to the desired rule, an estimate can be created. By creating instances that exactly match the available rules as well as for the input, a distance function can be used to find the instances, and by extension the rules, that most closely resemble the input. Based on these, an estimate can be calculated. The estimates may then be stored in the knowledge base to be compared with the observed metrics once the campaign has been run.

Since each rule has a confidence interval of succesful classification, multiple rules that are to be merged can be weighted based on their accuracy. For a single rule, the probability of success (i.e. user action on impression), disregarding the confidence interval, is calculated by taking the number of correctly classified instances, $\rho * N$, divided by the total number of matching instances N.

8 Verification

The result of the thesis work would be verified by applying the results to the problem suggested by Duego in a real-world setting. The suggested advertisement and target groups from the system will be added to their online marketing portfolio which will allow us to analyze the effectiveness of the software. This can be done continuously over the course of the development in addition to more traditional software unit and system testing.

The comparison metric to be used is the estimated click rate calculated for each rule set. This number can be compared with the actual click rate that the campaign which is created based on the rules receives. Formally, let $S_1 = (C, \hat{M})$ be the campaign C and its estimated metric $\hat{M}$, and $S_2 = (C, M)$ the actual outcome of the campaign. If the difference of the values is not statistically significant, i.e. $|\hat{M} - M| < \epsilon$, the suggestion is considered successful. If the difference is more notable, it might imply that the estimation function requires tuning.

9 Project plan

The proposed workflow of this thesis is inspired by the agile practices that perhaps most notably have become common in the software development industry. The most important methods that should be applied to this work are iterative project development and frequent "customer collaboration", in this case thesis suporvisors from both Chalmers University of Technology and Duego Technologies AB.

The project is expected to run for about 20 weeks. By dividing this relatively long period of time into iterations of four weeks, the project will be easier to manage from my point of view as well as provide ample opportunity for the supervisors to have their feedback incorporated.

The initial iterations will most likely require a heavy focus on research and studies of the related fields. The following iterations will include more work on the practical side of the projects. Once the research has progressed to an extent

that the problem space as well as the approach of the solution are well defined, more effort can be put into the development of the final application that is to be used by Duego. A preliminary list of milestones is listed below.

1. Finish extended proposal (June 19)

2. Demo a basic prototype of the system (July 9)

3. Full description and processing of input data (August 14)

4. Research should be finalized and draft of final paper submitted to supervisors (September 11)

5. Submit final paper to examinator and working system to Duego (October 9)

10   Thesis outline

- Abstract (0.5 pages)

- Introduction (0.5 pages)

- Problem description (1-2 pages)

- Related works (1 page)

- Solution (5-7 pages)

  - Analysis step
  - Planning step

- Discussion (1 page)

- Conclusion (0.5 pages)

References

Rakesh Agrawal, Tomasz Imielinski, and Arun Swami. Mining Association Rules between Sets of Items in Large Databases. In *SIGMOD Record*, volume 22, page 207, 1993.

N.V. Chawla, N. Japkowicz, and A. Kotcz. Editorial: special issue on learning from imbalanced data sets. *ACM SIGKDD Explorations Newsletter*, 6(1):1–6, 2004.

Ming-Syan Chen, Jiawei Han, and Philip S. Yu. Data Mining: An Overview from a Database Perspective. In *IEEE Transactions on Knowledge and Data Engineering*, volume 8, pages 866–883, 1996.

N. Chu, L. Ma, P. Liu, Y. Hu, and M. Zhou. A comparative analysis of methods for probability estimation tree. *WSEAS Transactions on Computers*, 10(3): 71–80, 2011.

Usama Fayyad, Gregory Piatetsky-shapiro, and Padhraic Smyth. Knowledge discovery and data mining: Towards a unifying framework. In *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD-96), Portland, Oregon, August 2-4, 1996*, pages 82–88. AAAI Press, 1996.

Eibe Frank and Ian H. Witten. Generating accurate rule sets without global optimization. In J. Shavlik, editor, *Fifteenth International Conference on Machine Learning*, pages 144–151. Morgan Kaufmann, 1998.

William J. Frawley, Gregory Piatetsky-Shapiro, and Christopher J. Matheus. Knowledge discovery in databases: An overview. *AI Magazine*, 13(3):57–70, 1992.

Stephen R. Garner. Weka: The waikato environment for knowledge analysis. In *Proceedings of the New Zealand computer science research students conference*, pages 57–64. Citeseer, 1995.

Michael Hahsler, Bettina Grün, and Kurt Hornik. Introduction to arules: Mining association rules and frequent item sets. *SIGKDD Explorations*, 2:4, 2007.

IBM. An architectural blueprint for autonomic computing. *Quality*, 36(June): 34, 2006.

N. Japkowicz and S. Stephen. The class imbalance problem: A systematic study. *Intelligent data analysis*, 6(5):429–449, 2002.

Marjan Mernik, Jan Heering, and Anthony M. Sloane. When and How to Develop Domain-Specific Languages. 37(4):316–344, 2005.

F. Provost and P. Domingos. Tree induction for probability-based ranking. *Machine Learning*, 52(3):199–215, 2003.

J. R. Quinlan. Induction of decision trees. *Mach. Learn.*, 1(1):81–106, March 1986. ISSN 0885-6125.

J.R. Quinlan. Simplifying decision trees. *International journal of man-machine studies*, 27(3):221–234, 1987a.

J.R. Quinlan. Generating production rules from decision trees. In *Proceedings of the Tenth International Joint Conference on Artificial Intelligence*, volume 30107, pages 304–307. Citeseer, 1987b.

M. Richardson, E. Dominowska, and R. Ragno. Predicting clicks: estimating the click-through rate for new ads. In *Proceedings of the 16th international conference on World Wide Web*, pages 521–530. ACM, 2007.

Ramakrishnan Srikant and Rakesh Agrawal. Mining quantitative association rules in large relational tables. *SIGMOD Rec.*, 25(2):1–12, 1996.

Ramakrishnan Srikant, Quoc Vu, and Rakesh Agrawal. Mining association rules with item constraints. In *Proceedings of the Third International Conference of Knowledge Discovery and Data Mining*, pages 67–73. AAAI Press, 1997.

Arie van Deursen, Paul Klint, and Joost Visser. Domain-specific languages: An annotated bibliography. 35(6):26–36, 2000.

Ian H. Witten, Eibe Frank, and Mark A. Hall. *Data Mining: Practical Machine Learning Tools and Techniques*. Morgan Kaufmann Series in Data Management Systems. Morgan Kaufmann, 3rd edition, 2011.